

Analysis of Algorithms

V. Adamchik

CSCI 570

Lecture 3

University of Southern California

Fall 2023

Heaps

Reading: chapter 3

Amortized Analysis

In a sequence of operations the worst case does not necessarily occur in each operation - some operations may take different times.

Therefore, a traditional worst-case per operation analysis can give overly pessimistic bound.

Consider insertions into an array

some operations take $O(n)$, others - $O(1)$

if the current array is full, the cost of insertion is linear;
if it is not full, insertion takes a constant time.

Therefore, amortized analysis is an alternative to the traditional worst-case analysis. Namely, we perform a worst-case analysis on a sequence of operations.

The Aggregate Method

The amortized cost of an operation is given by $\frac{T(n)}{n}$, where $T(n)$ is the upper bound on **the total cost** of n operations.

Example: unbounded array (with a doubling-up resizing policy)

Insertions: 1, 2, 3, 4, 5, 6, 7, 8, 9, ..., 2^{n+1}

Insertion Cost: 1, 1, 1, 1, 1, 1, 1, 1, ..., 1

Copy Cost: 0, 1, 2, 0, 4, 0, 0, 0, 8, ..., 2^n

In lecture 2 we computed the average cost per insert: $O(1)$

It is important to realize that we achieve a great amortized cost just because we have implemented a clever resizing policy!

Review Questions

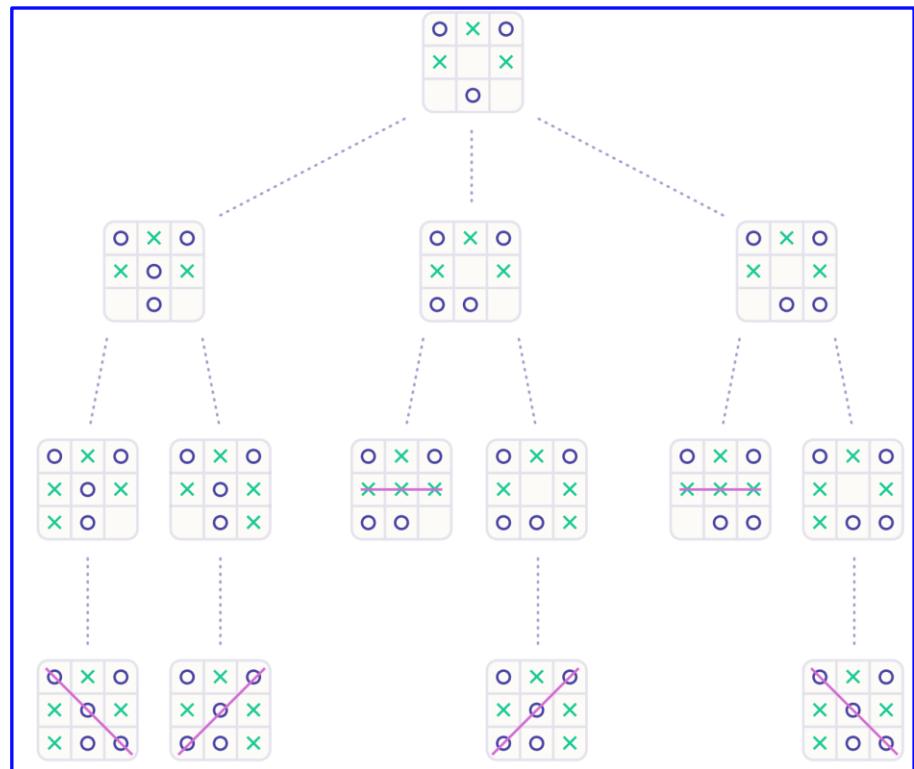
2. (T/F) Amortized analysis is used to determine the average runtime complexity of an algorithm.
3. (T/F) Compared to the worst-case analysis, amortized analysis provides a more accurate upper bound on the performance of an algorithm.
4. (T/F) The total amortized cost of a sequence of n operations gives a lower bound on the total actual cost of the sequence.
5. (T/F) Amortized constant time for a dynamic array is still guaranteed if we increase the array size by 5%.
6. (T/F) If an operation takes $O(1)$ expected time, then it takes $O(1)$ amortized time.
7. Suppose you have a data structure such that a sequence of n operations has an amortized cost of $O(n \log n)$. What could be the highest actual time of a single operation?
 $O(1), O(1), \dots, O(1), O(n^2) \rightarrow \text{Total} = O(n^2)$
 $O(1), O(1), \dots, O(1), O(n \log n)$ Ans: $O(n \log n)$

Heap and Priority Queue for Solving Optimization Problems

In this lecture, we will discuss a data structure that allows us to quickly access the highest priority element.

- ① sort
- ② Heap

Note: Never Use Heaps for searching · Heaps are used for solving optimization problems



To win a game you cannot simply run a DFS/BFS, among all possible moves you have to choose the best move!

BST

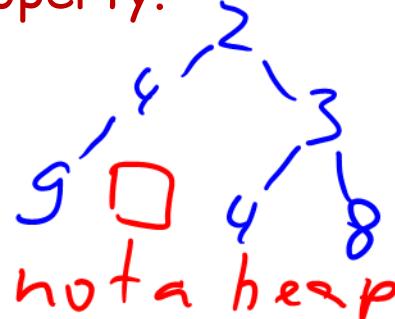
$L < P < R$

Binary min-Heap

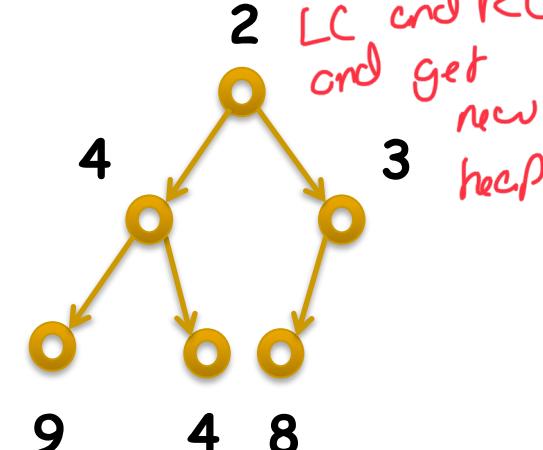
$L > P, R > P$

A binary heap is a **complete** binary tree which satisfies the **heap ordering property**.

1. Structure Property
2. Ordering Property



Binary Heap
is never unique
as we can swap
LC and RC
and get new
heap

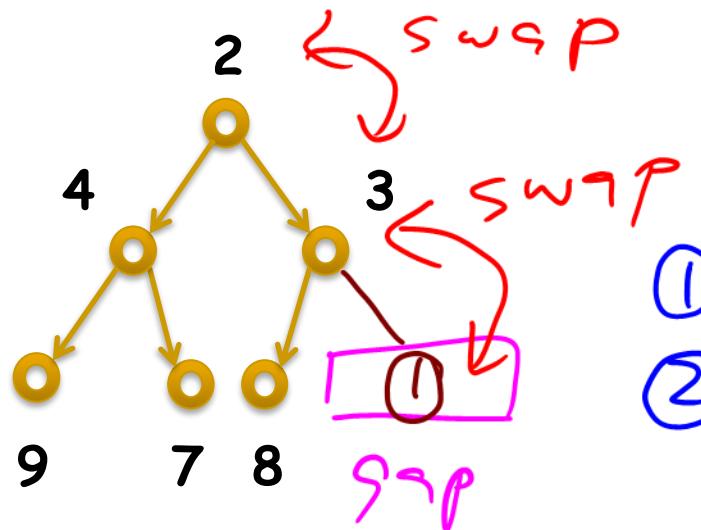


0	1	2	3	4	5	6	7
X	2	4	3	9	4	8	gap

Consider k -th element of the array,

- its left child is located at $2*k$ index
- its right child is located at $2*k+1$ index
- its parent is located at $k/2$ index

insert (tree reps)

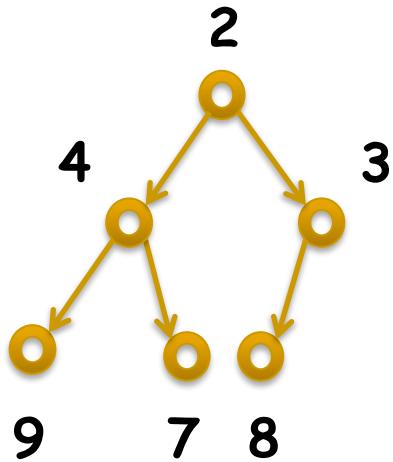


Algorithm:

- ① insert to a gap
- ② restore the ordering property

Runtime = # of swaps in the worst-case
 $O(\log n)$

insert (array reps)



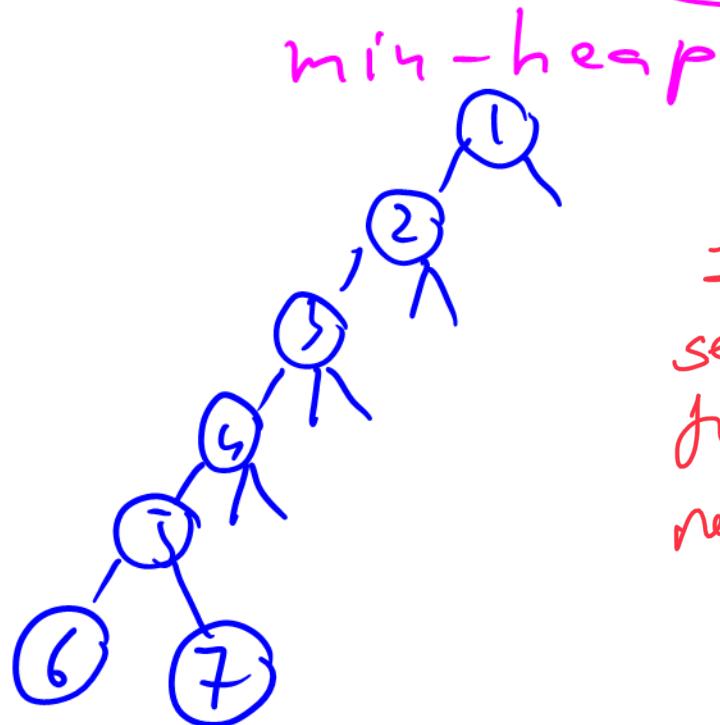
0	1	2	3	4	5	6	7
X	2	4	3	9	7	8	1
				1			3
	1	2					3

Implementation: a single for-loop
percolation

Discussion Problem 1

$$63 = 64 - 1 = 2^6 - 1$$

The values $1, 2, 3, \dots, 63$ are all inserted (in any order) into an initially empty min-heap. What is the smallest number that could be a leaf node?

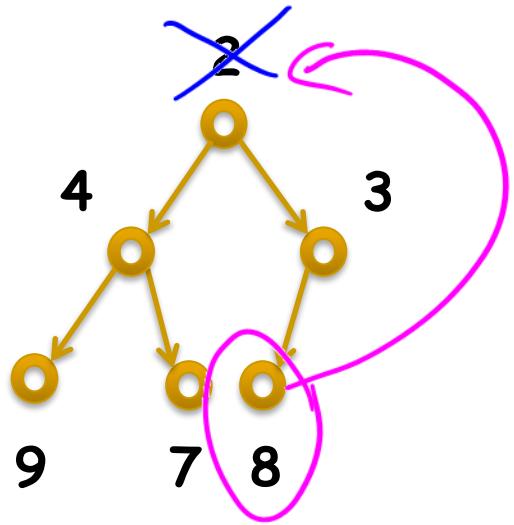


Proof by example.

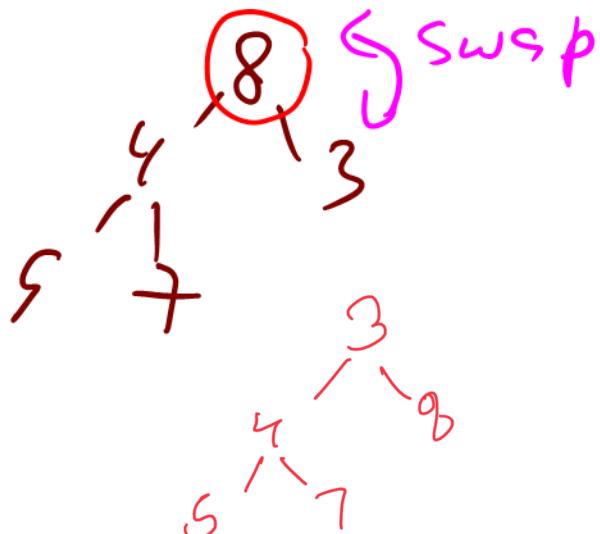
In this suppose we need to search 6 it will not take constant time to do it so sleep should never be used for searching

deleteMin (tree reps)

remove the root

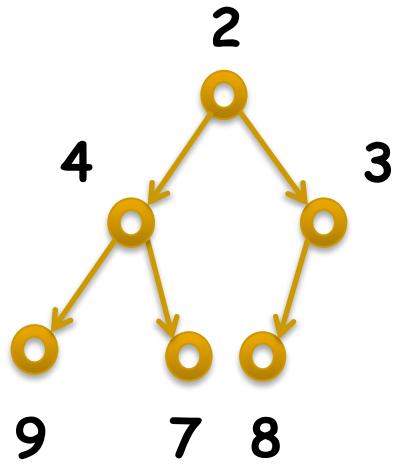


- ① preserve the structural property: by moving the last item to the root
- ② percolate down



Runtime: $O(\log n)$

deleteMin (array reps)



0	1	2	3	4	5	6	7
X	2	4	3	9	7	8	
X	8	4	3	9	7		
	3	8					

Implementation: a single for-loop

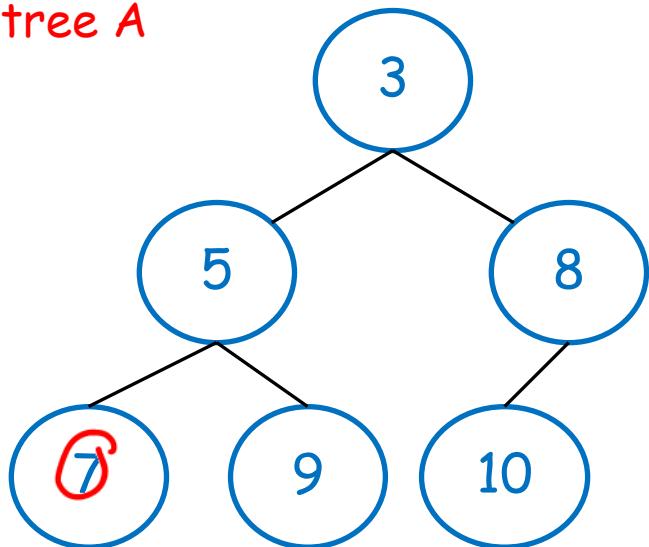
Discussion Problem 2

Suppose you have two binary min-heaps, A and B, with a total of n elements between them. You want to discover if A and B have a key in common. Give a solution to this problem that takes time $O(n \log n)$.

Do not use the fact that heaps are implemented as arrays, use only API operations: insert and deleteMin.

Runtim^e: $O(n \log n)$

tree A



Compare $3 < 6$

A. delete

$5 < 6$

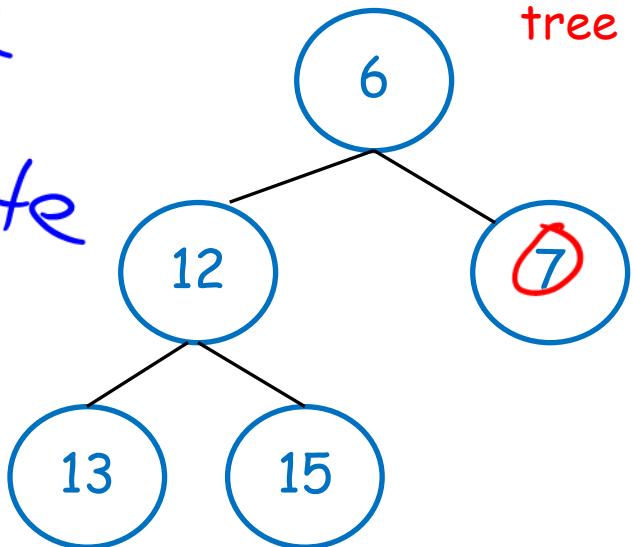
A. delete

$7 > 6$

B. delete

$7 = 7$

tree B

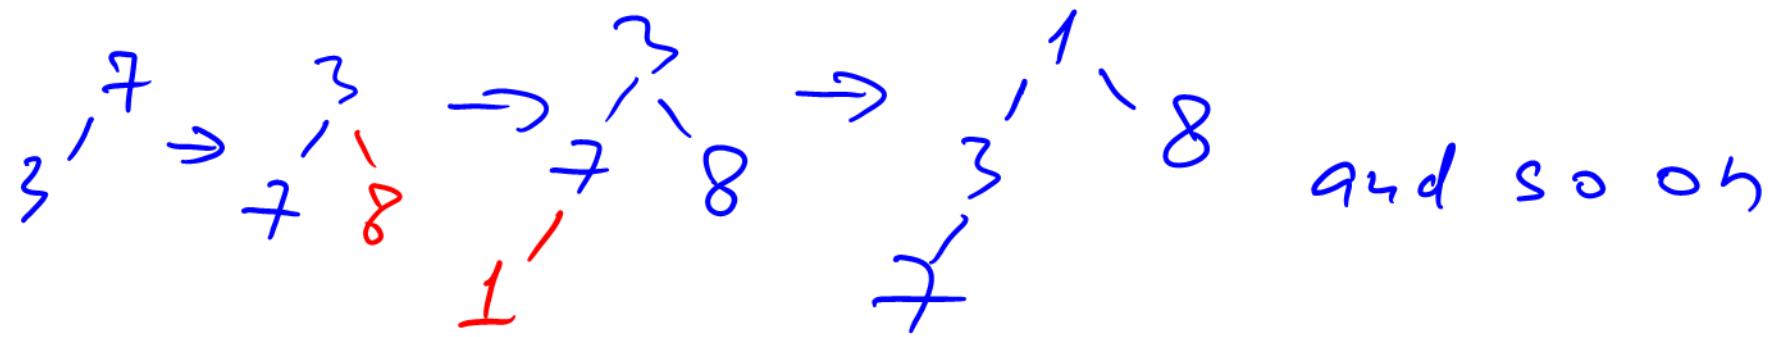


Build a BST. $O(n^2)$

Build a Heap by Insertion

Given an array - turn it into a heap.

insert 7, 3, 8, 1, 4, 9, 4, 10, 2, 0 into an initially empty heap.



Runtime: $O(h \cdot \log h)$

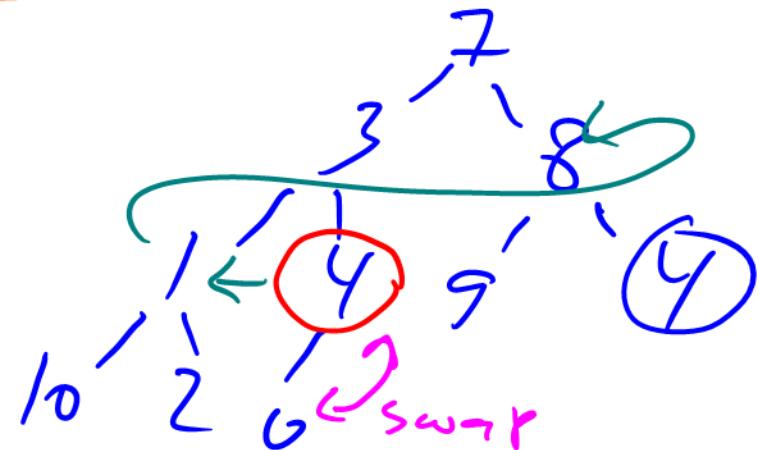
Total work: $\log 2 + \log 3 + \log 4 + \dots + \log h =$

$$\approx \log(h!) = O(h \log h)$$

Build a Heap in O(n)

Heapify:

7, 3, 8, 1, 4, 9, 4, 10, 2, 0



Algorithm:

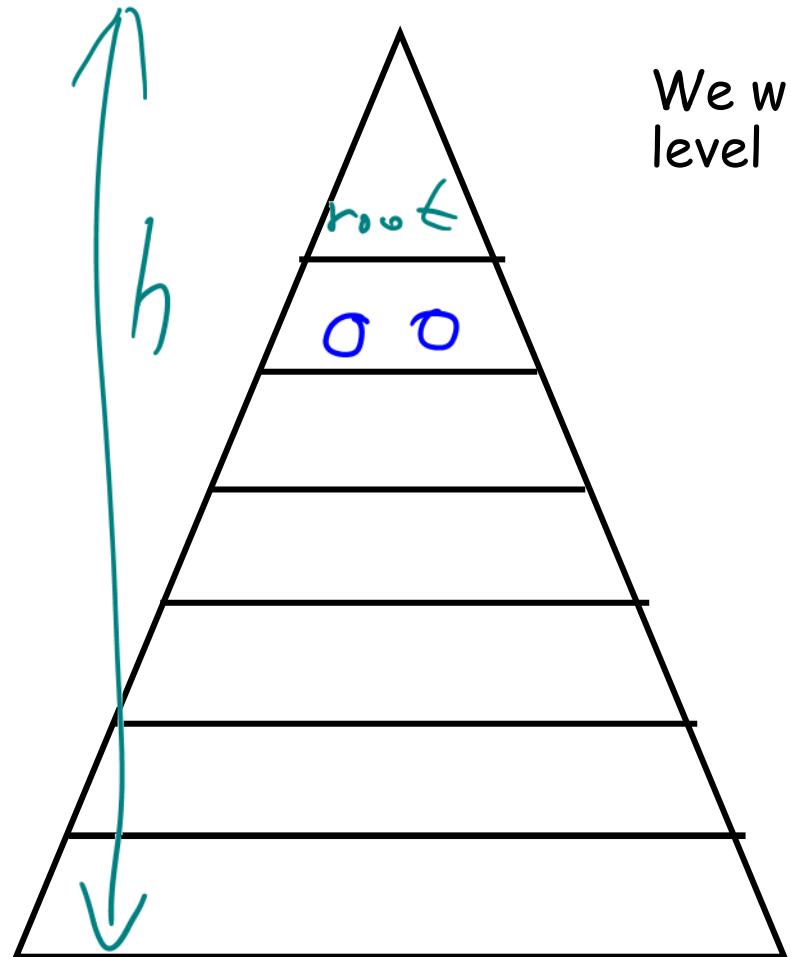
① start at index $n/2$

② compare with its children, and swap with the smallest child if it exists.

③ go to the left sibling

④ Repeat #2

Complexity of heapify



We will count the max number of swaps at each level

height	# of nodes	# of swaps
0 root	1	$\leftarrow h$
1	$2 \leftrightarrow h-1$	
2	$4 \leftrightarrow h-2$	
...
$h-1$	$2^{h-1} \leftrightarrow 1$	

Complexity of heapify

$$\begin{aligned} \text{Runtime} &= T_0 / n / \text{Cost} = h-1 \\ &= 1 \times h + 2 \times (h-1) + 4 \times (h-2) + \dots + 2^{h-1} \times 1 = \\ &= \sum_{k=0}^{h-1} 2^k \times (h-k) = \text{Wolfram Alpha} \\ &= O(n) \end{aligned}$$

Break, 5 mins

Discussion Problem 3

How would you sort using a binary heap?

What is its runtime complexity?

Sorting using BST: $O(n^2)$

① make a BST: $O(n^2)$

② traversal, in-order, $O(n)$

Heapsort algorithm:

① make a heap: $O(n)$

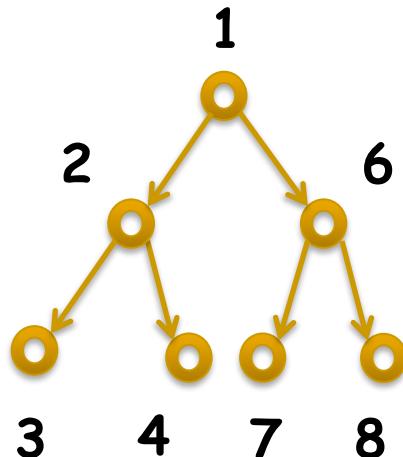
② delete: $O(n \log n)$

HEAPSORT

Run delMin n-times

$O(n \log n)$

in-place
nonstable



0	1	2	3	4	5	6	7
	1	2	6	3	4	7	8

	2	3	6	8	4	7	1
--	---	---	---	---	---	---	---

	3	4	6	8	7	2	1
--	---	---	---	---	---	---	---

	4	8	6	7	3	2	1
--	---	---	---	---	---	---	---

Discussion Problem 4

How would you merge two binary min-heaps?

What is its runtime complexity? $O(n)$

① offline algorithms

all data is available; you can preprocess
the data

$O(n)$

② online algorithms

(streaming data)

heapify

Runtime $O(n \log n)$ by insertions

Discussion Problem 5

Devise a heap-based algorithm that finds k largest elements out of n elements. Assume that $n > k$. What is its runtime complexity?

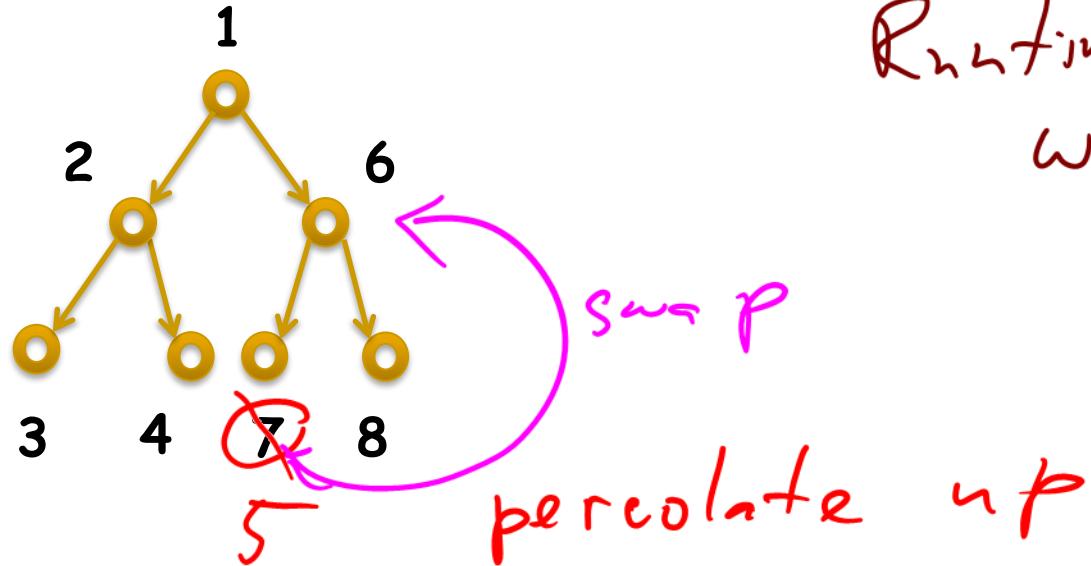
① offline (all data is available)

$O(n+k \log k)$ Algorithm:
a) build a max-heap
b) ran delete k times

② online (streaming data)

$O(K + (n-k) \log k)$ Algorithm:
a) wait for K items, then
 build a heap, min-heap
b) wait for $(K+1)^{st}$ item
c) new item \leq root, ignore it; o.w. run deleteMin , then insert

decreaseKey



Runtime: $O(\log n)$
worst-case

how do you find 7 in $O(1)$ time?

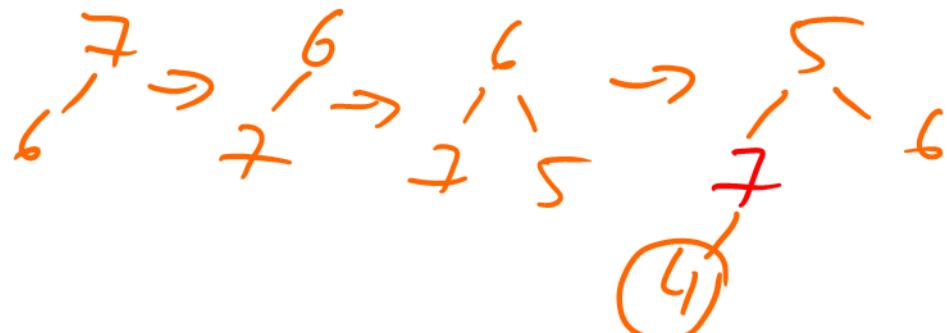


Add [1, 1, 1, 1] hashtable

A new kind of heaps

We want to create a heap with a better amortized complexity of insertion. This example will demonstrate that binary heaps do not provide a better upper bound for the worst-case complexity.

Insert $7, 6, 5, 4, 3, 2, 1$ into an empty binary min-heap.



The total number of swaps
 $\sum_{k=0}^{\log n - 1} k \cdot 2^k = O(n \log n)$

$$\begin{aligned} AC(\text{per insert}) &= \frac{O(n \log n)}{n} = \log n \end{aligned}$$

$$(1+x)^3 = \underbrace{1}_{} + \underbrace{3}_{}x + \underbrace{3}_{}x^2 + \underbrace{x^3}_{} + \dots$$

Not a
Complete
tree
1:n
Binom
tree

Binomial Trees B_k

The binomial tree B_k is defined as

1. B_0 is a single node
2. B_k is formed by joining two B_{k-1} trees

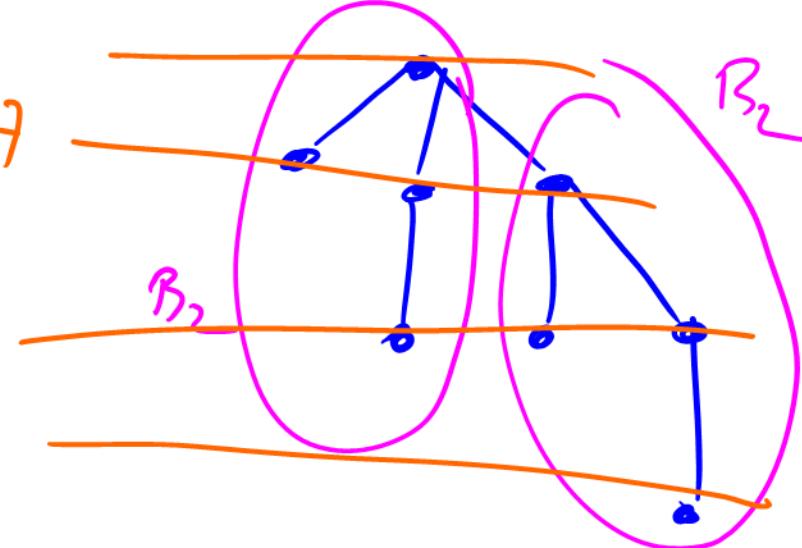
) but it follows
Binom tree ordering
 $P < \text{child}$

B_0

B_1

B_2

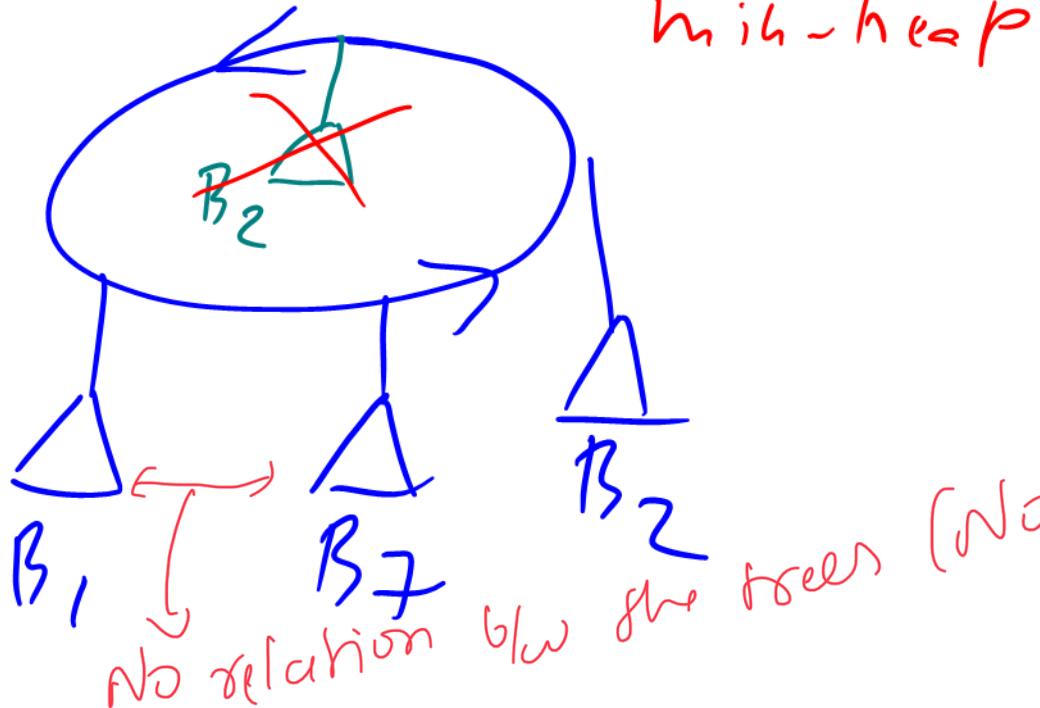
B_3



Binomial Heaps

→ Not just like Binomial heap
unique

A binomial heap is a collection (a linked list or a queue) of at most $\text{Ceiling}(\log n)$ binomial trees (of unique rank) in increasing order of size where each tree has a heap ordering property.



Discussion Problem 6

Given a sequence of numbers: 3, 5, 2, 8, 1, 5, 2, 7.

Draw a binomial heap by inserting the above numbers reading them from left to right

$$3 \leftrightarrow 5$$

B_0

$$\begin{matrix} 3 & \leftrightarrow & 2 & \leftrightarrow & 8 \\ | & & | & & | \\ 5 & & B_0 & & B_0 \\ & & \text{merge} & & \end{matrix}$$

$$\begin{matrix} 3 & \leftrightarrow & 2 \\ | & & | \\ 1 & & 8 \\ B_1 & & B_1 \\ \text{merge} & & \end{matrix} \quad \begin{matrix} 2 & / \\ 1 & \\ 3 & \\ 1 & \\ B_2 \end{matrix}$$

Discussion Problem 7

How many binomial trees does a binomial heap with 25 elements contain?

What are the ranks of those trees?

$$25_2 = \binom{10+8+1}{2} = 11001$$

$B_4 \quad B_3 \quad B_6$

Insertion

$O(\log n)$

What is its worst-case runtime complexity?

$$\begin{array}{r} 152 = 1111 \\ \times 41 \\ \hline 16000 \end{array}$$

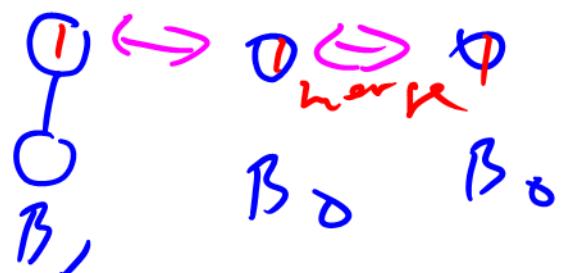
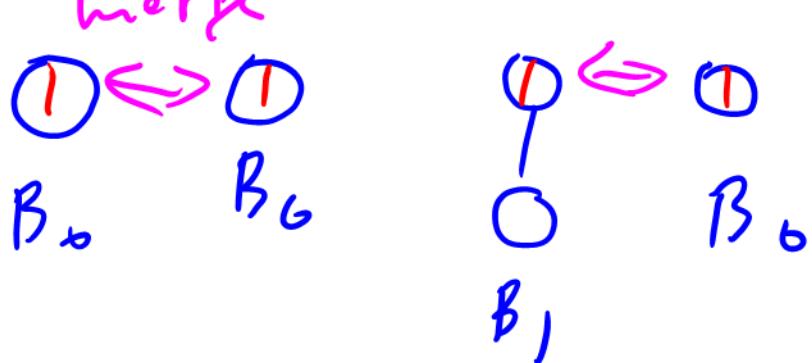
What is its amortized runtime complexity?

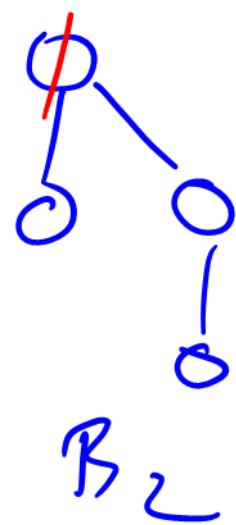
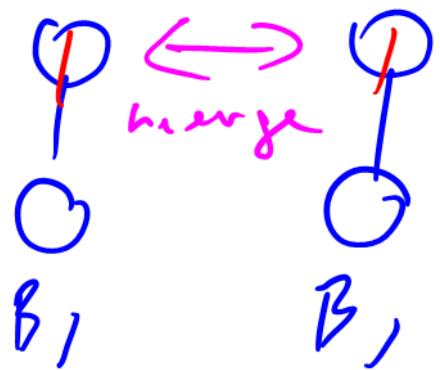
Use an accounting method.

Lecture 2

how many tokens
for each insert?

2 tokens





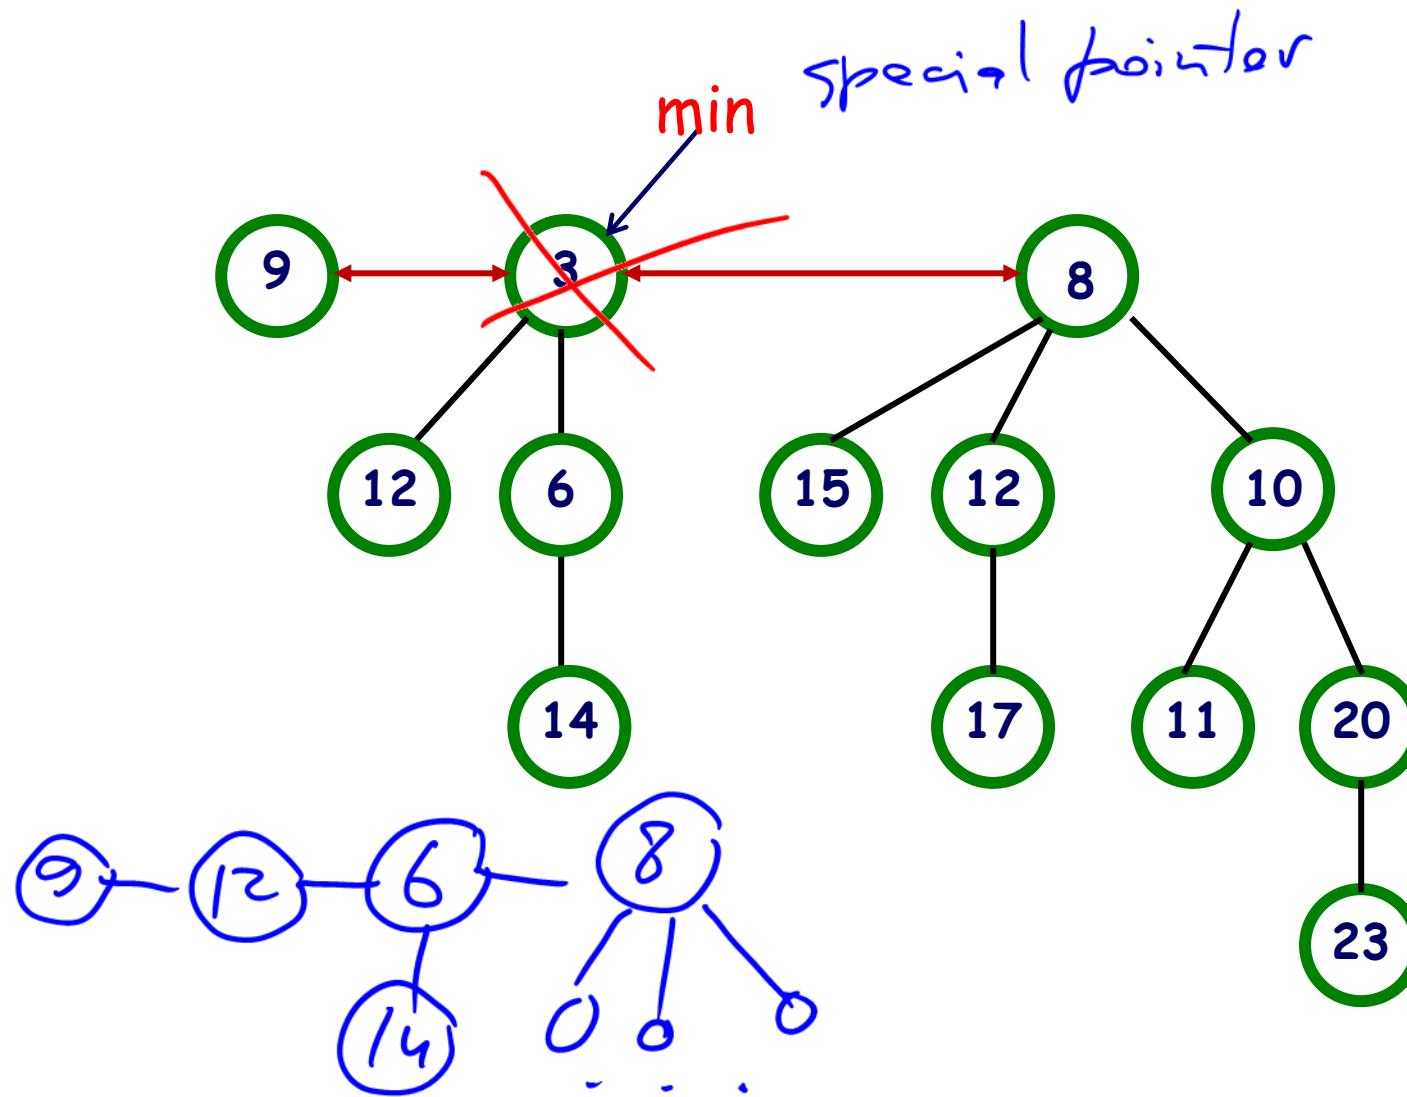
Building : Binomial vs Binary Heaps

The cost of inserting n elements into a **binary** heap, one after the other, is $\Theta(n \log n)$ in the worst-case. This is an online algorithm.

If n is known in advance (an offile algorithm), we run **heapify**, so a **binary** heap can be constructed in time $\Theta(n)$.

The cost of inserting n elements into a **binomial** heap, one after the other, is $\Theta(n)$ (amortized cost), even if n is not known in advance.

deleteMin()



deleteMin()

Algorithm:

- ① delete the min, $O(1)$
- ② move subtrees to the left (LL) level
- ③ traverse a LL and merge trees of the same rank

Runtime: $O(\log n)$

- ④ update the min-pointer, $O(\log n)$

Discussion Problem 8

Devise an algorithm for merging two binomial heaps and discuss its complexity. Merge $B_0B_1B_2B_4$ with B_1B_4 .

Merging two binary heap — $O(n)$
binomial — $O(1/\log n)$

Algorithm for binomial heaps

- ① merge two LL, $O(1)$
- ② traverse and merge binomial trees of the same rank, $O(1/\log n)$

$$\begin{array}{r} + 10111 \\ 10010 \\ \hline 101001 \end{array}$$

$B_5 B_3 B_0$

Heaps

"lazy"

	Binary	Binomial	<u>Fibonacci</u>
findMin	$\Theta(1)$	$\Theta(1)$	
deleteMin	$\Theta(\log n)$	$\Theta(\log n)$	
insert	$\Theta(\log n)$	$\Theta(1)$ (ac)	
decreaseKey	$\Theta(\log n)$	$\Theta(\log n)$	$O(1)$ ac
merge	$\Theta(n)$	$\Theta(\log n)$	

ac - amortized cost.

FIBONACCI HEAPS

Idea: relaxed (lazy) binomial heaps

Goal: decreaseKey in $O(1)$ ac.

The algorithm is outside of the scope of this course.

Heaps

	Binary	Binomial	Fibonacci
findMin	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
deleteMin	$\Theta(\log n)$	$\Theta(\log n)$	$O(\log n)$ (ac)
insert	$\Theta(\log n)$	$\Theta(1)$ (ac)	$\Theta(1)$
decreaseKey	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(1)$ (ac)
merge	$\Theta(n)$	$\Theta(\log n)$	$\Theta(1)$ (ac)